

scikit-learn

C. Travis Johnson

April 27, 2017

Mines Linux Users Group

Introduction

Machine Learning - What is it really?

- Goal: Extract Knowledge from Data
- Sometimes called predictive analysis or statistical learning
- Given a large matrix of observations X , fit a function $f(x)$ that maps observation x to a response variable y

Machine Learning - What is it really?

- Goal: Extract Knowledge from Data
- Sometimes called predictive analysis or statistical learning
- Given a large matrix of observations X , fit a function $f(x)$ that maps observation x to a response variable y

Machine Learning - What is it really?

- Goal: Extract Knowledge from Data
- Sometimes called predictive analysis or statistical learning
- Given a large matrix of observations X , fit a function $f(x)$ that maps observation x to a response variable y

Important Terms

Classifiers Algorithms that learn functions to map observations to a *discrete* response. E.g., is this tumor malignant or benign? Is this email spam or not?

Regressors Algorithms that learn functions to map observations to a *continuous* response. E.g., how much should this house cost?

Underfitting The learned function is too simple. “We barely studied for the exam.”

Overfitting The learned function is too complex. “We memorized all the practice problems, but don’t understand the material.”

Generalization How well does the learned function extend to new observations?

Scikit-Learn: Machine Learning in Python

- Provides many machine learning tools with a common Estimator interface¹
- Built in helpers for common ML tasks (e.g., metrics, preprocessing)
- Easily combine algorithms to make a complex pipeline²
- Relies heavily on `numpy` and `scipy`, often used with `pandas`

¹<http://scikit-learn.org/stable/developers/contributing.html#apis-of-scikit-learn-objects>

²Sound familiar?

Scikit-Learn: Machine Learning in Python

- Provides many machine learning tools with a common Estimator interface¹
- Built in helpers for common ML tasks (e.g., metrics, preprocessing)
- Easily combine algorithms to make a complex pipeline²
- Relies heavily on numpy and scipy, often used with pandas

¹<http://scikit-learn.org/stable/developers/contributing.html#apis-of-scikit-learn-objects>

²Sound familiar?

Scikit-Learn: Machine Learning in Python

- Provides many machine learning tools with a common Estimator interface¹
- Built in helpers for common ML tasks (e.g., metrics, preprocessing)
- Easily combine algorithms to make a complex pipeline²
- Relies heavily on `numpy` and `scipy`, often used with `pandas`

¹<http://scikit-learn.org/stable/developers/contributing.html#apis-of-scikit-learn-objects>

²Sound familiar?

Scikit-Learn: Machine Learning in Python

- Provides many machine learning tools with a common Estimator interface¹
- Built in helpers for common ML tasks (e.g., metrics, preprocessing)
- Easily combine algorithms to make a complex pipeline²
- Relies heavily on numpy and scipy, often used with pandas

¹<http://scikit-learn.org/stable/developers/contributing.html#apis-of-scikit-learn-objects>

²Sound familiar?

Supervised Learning

Learning to Predict Breast Cancer

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split

cancer = load_breast_cancer()      # Get some data
X_train, X_test, y_train, y_test = train_test_split(
    cancer.data, cancer.target,
    stratify=cancer.target, random_state=1337)

tree = DecisionTreeClassifier(random_state=7331)
tree.fit(X_train, y_train)      # Learn a Decision Function
```

Evaluating Accuracy of a Model

```
# How well did we do?  
train_acc = tree.score(X_train, y_train)  
test_acc = tree.score(X_test, y_test)  
print("Training Accuracy: {:.3f}".format(train_acc))  
print("Testing Accuracy: {:.3f}".format(test_acc))  
# Training Accuracy: 1.000  
# Testing Accuracy: 0.923
```

Other Supervised Learning Models

- Decision trees are a common first step, because they're easy to interpret and don't require much preprocessing
- Decision trees are prone to overfitting, so a good improvement is the `RandomForest`
- Support Vector Machines, Logistic/Linear Regression, and Artificial Neural Networks are commonly the first algorithms studied
- See the `scikit-learn` documentation for a comprehensive guide of available algorithms

Other Supervised Learning Models

- Decision trees are a common first step, because they're easy to interpret and don't require much preprocessing
- Decision trees are prone to overfitting, so a good improvement is the `RandomForest`
- Support Vector Machines, Logistic/Linear Regression, and Artificial Neural Networks are commonly the first algorithms studied
- See the `scikit-learn` documentation for a comprehensive guide of available algorithms

Other Supervised Learning Models

- Decision trees are a common first step, because they're easy to interpret and don't require much preprocessing
- Decision trees are prone to overfitting, so a good improvement is the `RandomForest`
- Support Vector Machines, Logistic/Linear Regression, and Artificial Neural Networks are commonly the first algorithms studied
- See the `scikit-learn` documentation for a comprehensive guide of available algorithms

Other Supervised Learning Models

- Decision trees are a common first step, because they're easy to interpret and don't require much preprocessing
- Decision trees are prone to overfitting, so a good improvement is the `RandomForest`
- Support Vector Machines, Logistic/Linear Regression, and Artificial Neural Networks are commonly the first algorithms studied
- See the `scikit-learn` documentation for a comprehensive guide of available algorithms

Becoming a “Data Scientist”

1. Get some (more) data
2. Pick an algorithm (or algorithm chain)
3. Train the model
4. Test generalization ability of trained model
5. Good enough? Done. Else, go back to step 1 or 2.

Then, tell people you're a genius . . . it's that easy!

Unsupervised Learning

Distinction from Supervised Learning

Supervised Learning You tell the model what the correct answers are for training examples.

Unsupervised Learning You ask the model to extract information from a dataset.

Unsupervised Clustering Partition data into similar groups.
Example: K-Means Clustering

Unsupervised Transformations Create new representations of data. Example: Principal Component Analysis

Model Evaluation and Improvement

Choice of Evaluation Metric

- Accuracy is not always the best metric for your system
- Plenty of others exist, pick the best for your business costs
- Look in the `sklearn.metrics` module for alternatives
- You can also use your own evaluation function!

Choice of Evaluation Metric

- Accuracy is not always the best metric for your system
- Plenty of others exist, pick the best for your business costs
- Look in the `sklearn.metrics` module for alternatives
- You can also use your own evaluation function!

Choice of Evaluation Metric

- Accuracy is not always the best metric for your system
- Plenty of others exist, pick the best for your business costs
- Look in the `sklearn.metrics` module for alternatives
- You can also use your own evaluation function!

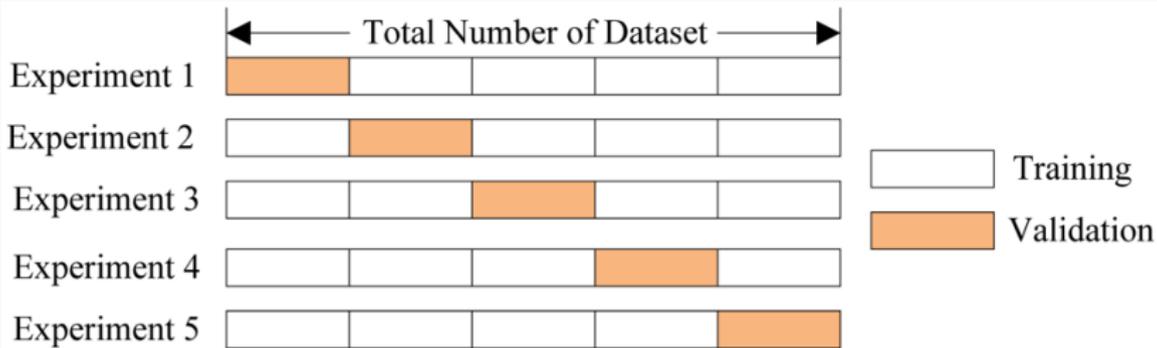
Choice of Evaluation Metric

- Accuracy is not always the best metric for your system
- Plenty of others exist, pick the best for your business costs
- Look in the `sklearn.metrics` module for alternatives
- You can also use your own evaluation function!

Cross Validation

Never Fit Models to Test Data! Ever!

Learning the parameters of a prediction function and testing it on the same data is a methodological mistake: a model that would just repeat the labels of the samples that it has just seen would have a perfect score but would fail to predict anything useful on yet-unseen data. This situation is called overfitting.



Grid Search with Cross Validation

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV

cancer = load_breast_cancer()    # Get some data
X_train, X_test, y_train, y_test = train_test_split(
    cancer.data, cancer.target,
    stratify=cancer.target, random_state=1337)

tree = DecisionTreeClassifier(random_state=7331)
search_grid = {'criterion': ['gini', 'entropy'],
               'max_depth' : [5, 10, 15, 20]}

# search_grid could also be a list of dicts
search = GridSearchCV(tree, search_grid, cv=5)
search.fit(X_train, y_train)
print(search.best_params_)
```

Pipelines

Use `Pipeline` to combine multiple estimators into a single estimator. Two conveniences:

1. Convenience: You only have to call `fit` and `predict` once on your data to fit a whole sequence of estimators.
2. Joint parameter selection: You can grid search over parameters of all estimators in the pipeline at once.

A Simple Pipeline

```
>>> from sklearn.pipeline import Pipeline
>>> from sklearn.svm import SVC
>>> from sklearn.decomposition import PCA
>>> estimators = [('reduce_dim', PCA()), ('clf', SVC())]
>>> pipe = Pipeline(estimators)
>>> pipe
Pipeline(steps=[('reduce_dim', PCA(copy=True, iterated_power='auto',
n_components=None, random_state=None, svd_solver='auto', tol=0.0,
whiten=False)), ('clf', SVC(C=1.0, cache_size=200, class_weight=None,
coef0=0.0, decision_function_shape=None, degree=3, gamma='auto',
kernel='rbf', max_iter=-1, probability=False, random_state=None,
shrinking=True, tol=0.001, verbose=False))])
```

Grid Search - Tuning a Complex Pipeline

```
from sklearn.pipeline import make_pipeline
from sklearn.svm import SVC
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
pipe = make_pipeline(PCA(), StandardScaler(), SVC())
params = dict(pca__n_components=[2, 5, 10],
              svc__C=[0.1, 10, 100])
grid = GridSearchCV(pipe, param_grid=params)
# Next, call grid.fit on some training data
# This will use cross validation to estimation performance using each
# combination of parameters for pipeline in params dict

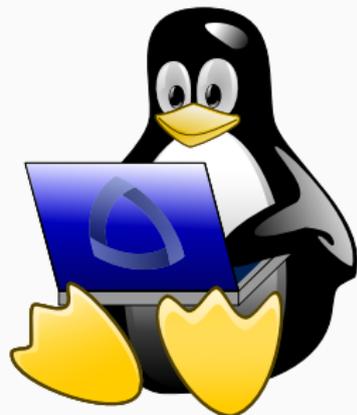
# With fitted model
print(grid.best_params_)
```

Questions?

Copyright Notice

This presentation was from the **Mines Linux Users Group**. A mostly-complete archive of our presentations can be found online at <https://lug.mines.edu>.

Individual authors may have certain copyright or licensing restrictions on their presentations. Please be certain to contact the original author to obtain permission to reuse or distribute these slides.



Colorado School of Mines
Linux Users Group